


```

File Edit View Search Terminal Help
Time taken: 0.055 seconds
hive> insert into flight values(123,2009,"pune",305);
FAILED: SemanticException [Error 10001]: Line 1:12 Table not found 'flight'
hive> insert into flight1 values(123,2009,"pune",305);
Query ID = cloudera_20240507073737_a52da88c-7e28-460b-9515-436bbf7017dc
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714984172459_0014, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1714984172459_0014/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1714984172459_0014
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-05-07 07:37:12,387 Stage-1 map = 0%, reduce = 0%
2024-05-07 07:37:17,774 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 0.99 sec
MapReduce Total cumulative CPU time: 990 msec
Ended Job = job_1714984172459_0014
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/soham.db/flight1/.hive-staging_hive_2024-05-07_07-37-06_135_8023104256826965805-1/-ext-10000
Loading data to table soham.flight1
Table soham.flight1 stats: [numFiles=1, numRows=1, totalSize=18, rawDataSize=17]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 0.99 sec HDFS Read: 4099 HDFS Write: 87 SUCCESS
Total MapReduce CPU Time Spent: 990 msec
OK
Time taken: 12.892 seconds
hive> select * from flight1;
OK
123 2009 pune 305
Time taken: 0.072 seconds, Fetched: 1 row(s)
hive> load data local inpath "f1.txt"
> overwrite into table flight1;
Loading data to table soham.flight1
Table soham.flight1 stats: [numFiles=1, numRows=0, totalSize=79, rawDataSize=0]
OK
Time taken: 0.243 seconds
hive> select * from flight1;
OK
123 2009 goa 306
124 2009 kolhapur 306
125 2009 mumbai 306
126 2009 kerala 306
Time taken: 0.039 seconds, Fetched: 4 row(s)
hive> create table nflight(fno int ,year int,source varchar(10));
OK
Time taken: 0.037 seconds
hive> drop table nflight;
OK

```

```

File Edit View Search Terminal Help
125 2009 mumbai 306
126 2009 kerala 306
Time taken: 0.039 seconds, Fetched: 4 row(s)
hive> create table nflight(fno int ,year int,source varchar(10));
OK
Time taken: 0.037 seconds
hive> drop table nflight;
OK
Time taken: 0.062 seconds
hive> create table nflight(fno int ,year int,source varchar(10))
> row format delimited
> fields terminated by ','
> lines terminated by '\n'
> stored as textfile;
OK
Time taken: 0.051 seconds
hive> insert into nflight values(123,2009,pune);
FAILED: SemanticException [Error 10293]: Unable to create temp file for insert values Expression of type TOK_TABLE_OR_COL not supported in insert/values
hive> insert into nflight values(123,2009,"pune");
Query ID = cloudera_20240507074444_aa2e5039-1a48-40c1-9b70-945c11b4fce7
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714984172459_0015, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1714984172459_0015/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1714984172459_0015
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
2024-05-07 07:44:54,880 Stage-1 map = 0%, reduce = 0%
2024-05-07 07:45:00,141 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.02 sec
MapReduce Total cumulative CPU time: 1 seconds 20 msec
Ended Job = job_1714984172459_0015
Stage-4 is selected by condition resolver.
Stage-3 is filtered out by condition resolver.
Stage-5 is filtered out by condition resolver.
Moving data to: hdfs://quickstart.cloudera:8020/user/hive/warehouse/soham.db/nflight/.hive-staging_hive_2024-05-07_07-44-48_371_1897369824036507295-1/-ext-10000
Loading data to table soham.nflight
Table soham.nflight stats: [numFiles=1, numRows=1, totalSize=14, rawDataSize=13]
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Cumulative CPU: 1.02 sec HDFS Read: 3894 HDFS Write: 83 SUCCESS
Total MapReduce CPU Time Spent: 1 seconds 20 msec
OK
Time taken: 12.964 seconds
hive> insert into nflight values(124,2009,"mumbai");
Query ID = cloudera_20240507074545_24328277-ca15-42ed-bb15-fd2d2719f123
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714984172459_0016, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1714984172459_0016/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1714984172459_0016

```

File Edit View Search Terminal Help

 [clou...
 [kom...
 [clou...
 [clou...
 [clou...
 [clou...
 [clou...
 [clou...
 clou...


```

Execution log at: /tmp/cloudera/cloudera_20240507074848_c0645944-52b2-400d-adf2-5cb38fa213ab.log
2024-05-07 07:48:28 Starting to launch local task to process map join; maximum memory = 1013645312
2024-05-07 07:48:29 Dump the side-table for tag: 1 with group count: 2 into file: file:/tmp/cloudera/180fe35e-ceaa-40c1-adf9-ec0d197cef96/hive_2024-05-07_07-48-25_767_231124971931338816-1/-local-10003/HashTable-Stage-3/MapJo
001--hashtable
2024-05-07 07:48:29 Uploaded 1 File to: file:/tmp/cloudera/180fe35e-ceaa-40c1-adf9-ec0d197cef96/hive_2024-05-07_07-48-25_767_231124971931338816-1/-local-10003/HashTable-Stage-3/MapJo
2024-05-07 07:48:29 End of local task; Time Taken: 1.076 sec.
Execution completed successfully
MapredLocal task succeeded
Launching Job 1 out of 1
Number of reduce tasks is set to 0 since there's no reduce operator
Starting Job = job_1714984172459_0017, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1714984172459_0017/
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1714984172459_0017
Hadoop job information for Stage-3: number of mappers: 1; number of reducers: 0
2024-05-07 07:48:35,957 Stage-3 map = 0%, reduce = 0%
2024-05-07 07:48:41,185 Stage-3 map = 100%, reduce = 0%, Cumulative CPU 0.92 sec
MapReduce Total cumulative CPU time: 920 msec
Ended Job = job_1714984172459_0017
MapReduce Jobs Launched:
Stage-Stage-3: Map: 1 Cumulative CPU: 0.92 sec HDFS Read: 6490 HDFS Write: 38 SUCCESS
Total MapReduce CPU Time Spent: 920 msec
OK
123 2009 306 pune
124 2009 306 mumbai
Time taken: 16.466 seconds, Fetched: 2 row(s)
hive> select * from flight1;
OK
123 2009 goa 306
124 2009 kolhapur 306
125 2009 mumbai 306
126 2009 kerala 306
Time taken: 0.038 seconds, Fetched: 4 row(s)
hive> create index flight_index on table flight1(fno)
> as 'org.apache.hadoop.hive.ql.index.compact.CompactIndexHandler'
> WITH DEFERRED REBUILD;
OK
Time taken: 0.166 seconds
hive> show tables;
OK
air flight
flight1
nflight
soham flight1 flight_index
values tmp table_1
values tmp table_2
values tmp table_4
values tmp table_5
Time taken: 0.008 seconds, Fetched: 8 row(s)
hive> [cloudera@quickstart ~]$

```